

# HW

## ① Serializability & 2 PL

	T <sub>1</sub>	T <sub>2</sub>	T <sub>3</sub>
t <sub>1</sub>	read(x)		
t <sub>2</sub>		read(y)	
t <sub>3</sub>			Write(z)
t <sub>4</sub>		read(x)	
t <sub>5</sub>	read(y)		

read(x) → In T<sub>1</sub>, no subsequent writer to x, no new edge.  
 read(y) → In T<sub>2</sub>, no subsequent writer to y, no new edge.  
 write(z) → x is subsequent read by T<sub>2</sub>. new edge T<sub>2</sub> → T<sub>3</sub>  
 read(x) → In T<sub>2</sub>, no subsequent writers to x, no new edge.

Serializable → no cycles

	T <sub>1</sub>	T <sub>2</sub>	T <sub>3</sub>
t <sub>1</sub>	lock(x)		
t <sub>2</sub>	read(x)		
t <sub>3</sub>	lock(y)		
t <sub>4</sub>	read(y)		
t <sub>5</sub>	write(x)		
t <sub>6</sub>	x-lock(x)	lock(y)	
	s.r(x)	read(y)	
		x-lock(y)	
		write(y)	
		x.r(y)	
		s.r(y)	

read(x) → No future write to x  
 read(y) → y is writer in T<sub>3</sub>: T<sub>1</sub> → T<sub>3</sub>  
 write(x) → no future write to x  
 read(y) → y is written in T<sub>3</sub>: T<sub>2</sub> → T<sub>3</sub>  
 write(y) → y is read in T<sub>2</sub>: T<sub>3</sub> → T<sub>2</sub>

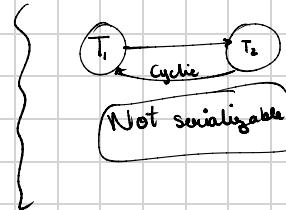
Not Serializable, procedure has cycles.

Follows 2PL

1.3

	T <sub>1</sub>	T <sub>2</sub>
t <sub>1</sub>	x.lock(x)	
t <sub>2</sub>	write(x)	
t <sub>3</sub>	x.r(x)	
t <sub>4</sub>	x.lock(x)	lock(y)
t <sub>5</sub>	read(y)	
t <sub>6</sub>	s.r(x)	

Doesn't follow 2PL because T<sub>1</sub> tries to get an x-lock on x after releasing it.



1.4

	T <sub>1</sub>	T <sub>2</sub>	T <sub>3</sub>
t <sub>1</sub>	lock(x)		
t <sub>2</sub>	read(x)		
t <sub>3</sub>		x.lock(y)	
t <sub>4</sub>		write(y)	
t <sub>5</sub>	write(x)		
t <sub>6</sub>			lock(x)
			read(x)

t<sub>1</sub>: read(x) → x has future write in T<sub>2</sub>: T<sub>1</sub> → T<sub>2</sub>  
 t<sub>2</sub>: write(x) → x has future read in T<sub>3</sub>: T<sub>2</sub> → T<sub>3</sub>  
 t<sub>3</sub>: write(y) → y has future read in T<sub>3</sub>: T<sub>1</sub> → T<sub>3</sub>  
 t<sub>4</sub>: x doesn't have future write elsewhere.

No Cycles = Serializable.

## 2) Degree of consistency:

T1 has a max degree of consistency of 0.

T2 has a max degree of consistency of 3.

## 3) Read Uncommitted:

- \* Isolation is allowing transactions to read uncommitted data from other transaction potentially leading to dirty reads.
- \* Equivalent degree of consistency: Degree 1: this is cause Read uncommitted allows the transaction to read data that has been read but not committed by the concurrent transaction.

## Read Committed:

- \* It ensures the transaction only reads data that has been committed by other transactions. It ensures that a transaction sees only consistent data & prevents any dirty reads, so it doesn't read uncommitted data.
- \* Equivalent degree of consistency: Degree 2 - it aligns w/ the principle that transactions don't commit any writes until the end of the transaction, ensuring consistency & preventing partial updates from being visible to other transactions.

## Repeatable Read:

- \* Ensuring that once a transaction reads a piece of data, it will see the same data for the remainder of that transaction, even if other transactions are modified or inserted data.
- \* Equivalent degree of consistency: Degree 0  $\rightarrow$  doesn't guarantee dirty data of other facts. Satisfying this because this degree prevents other transactions from modifying data reads by the current transaction.

## S Serializable

- \* This prevents all isolations. Creates a shared lock on all statements. This follows a degree of consistency of 3. This degree satisfies by ensuring that data read by current transaction remains consistent until the transaction completes.

# 4) Recovery & Freies:

## a) Action 0:

- \* Action: T1 updates P7, P7 brought into buffer
- \* XT: {T1: 00}
- \* DPT: {P7: 00}

## Action 1:

- \* Action: T0 updates P9, P9 brought into buffer, P9 flushed to disk
- \* XT: {T1: 00, T0: 10}
- \* DPT: {P7: 00}

## Action 2: LSN = 20

- \* Action: T1 updates P8, P8 brought into buffer, P8 flushed to disk
- \* XT: {T1: 20, T0: 10}
- \* DPT: {P7: 00}

## Action 3: LSN=30

- \* Action: begin checkpoint
- \* No change to XT & DPT

## Action 4: LSN=40

- \* Action: end checkpoint
- \* No change to XT & DPT

## Action 5: LSN=50

- \* Action: T1 updates P9
- \* XT {T1: 50, T0: 10}
- \* DPT {P7: 00, P9: 50}

## Action 6: LSN=60

- \* Action: T2 updates P6
- \* XT {T1: 50, T0: 10, T2: 60}
- \* DPT {P7: 00, P9: 50, P6: 60}

Action 7: LSN=70

- \* Action: T1 updates P5
- \* XT { T1: 70, T0: 10, T2: 60 }
- \* DPT { P7: 00, P9: 50, P6: 60, P5: 70 }

Action 8: LSN=80

- \* Action: T1 updates P7, P6 flushed to disk
- \* XT { T1: 80, T0: 10, T2: 60 }
- \* DPT { P7: 00, P9: 50, P5: 70 }

Action 9:

Restart

trans ID	recLSN	DPT
XT	T1: 80 T0: 10 T2: 60	

Page ID	last LSN
	PT: 80

#### 4B) The analysis

- \* It begins where the checkpoint ends
- \* the analysis in this case starts at log time 4 with LSN=40

Time	LSW	Transition Table	Dirty Page Table
4	40	T1: 20, T0: 10	P7: 00
5	50	T1: 50, T0: 10	P7: 00, P9: 50
6	60	T1: 50, T0: 10, T2: 60	P7: 00, P9: 50, P6: 60
7	70	T1: 70, T0: 10, T2: 60	P7: 00, P9: 50, P6: 60, P5: 70
8	80	T1: 80, T0: 10, T2: 60	P7: 00, P9: 50, P5: 70
9	Restart	Restart	Restart

\* Analysis is starting here

4c) The redo phase starts at 00, that is because that is the earliest page LSN from the end checkpoint.

Time	LSN	Log	Redo or Not
0	00	update: T1 updates P7	Redo
1	10	update: T0 updates P9	NOT
2	20	T1 updates P8	Not
3	30	begin checkpoint	Not
4	40	end checkpoint	Not
5	50	T1 updates P9	Redo
6	60	T2 updates P6	Not
7	70	T1 updates P5	Redo
8	80	T1 updates P7	Redo

## 5) Recovery and Aries

The ARIES and InnoDB Crash Recovery algorithms share striking similarities. They both encompass “Analysis” and “Redo” phases, which are crucial for data restoration. While ARIES maintains both an undo and redo log, InnoDB primarily relies on a redo log. Additionally, both algorithms address incomplete transactions through an undo mechanism. Notably, InnoDB offers a unique feature where, if all changes preceding the crash have been successfully flushed, it skips the redo phase altogether.